

Towards GPU-accelerated multimessenger inference of neutron star mergers and dense matter physics

Thibeau Wouters



Utrecht
University

Nikhef

Table of Contents

① Introduction

② Methods

③ Applications

④ Neural priors

⑤ Outlook and conclusion

Contents

① Introduction

Neutron stars and equation of state

Multimessenger astrophysics

Future detectors and challenges

② Methods

③ Applications

④ Neural priors

⑤ Outlook and conclusion

Contents

① Introduction

Neutron stars and equation of state

Multimessenger astrophysics

Future detectors and challenges

② Methods

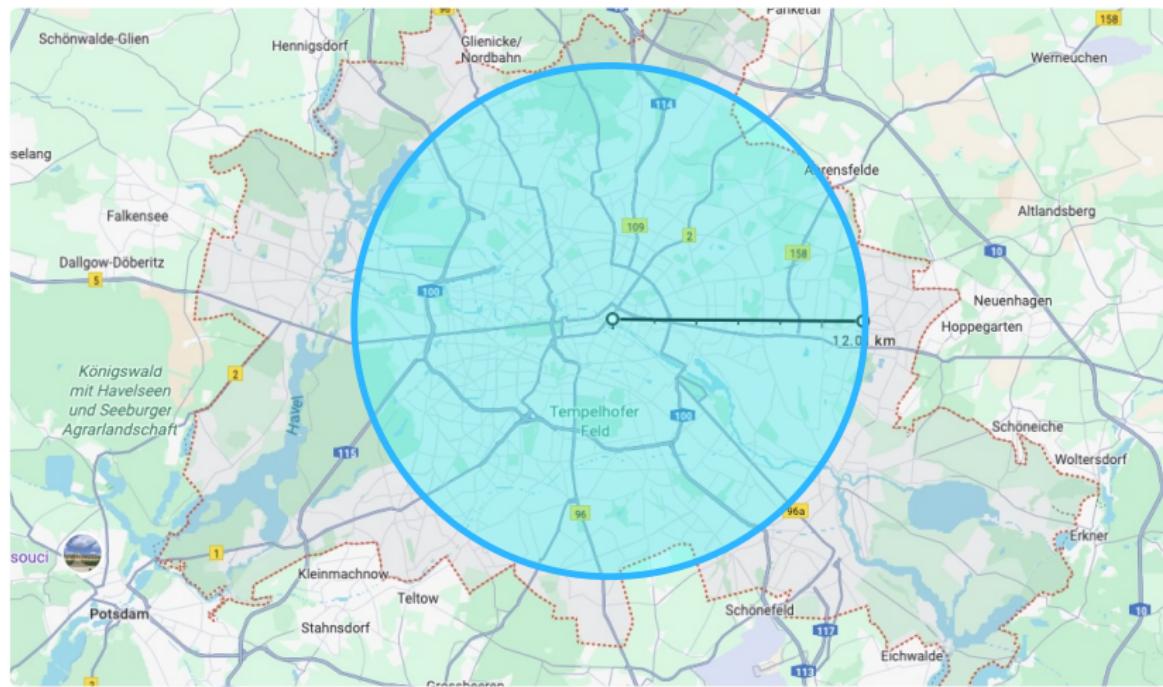
③ Applications

④ Neural priors

⑤ Outlook and conclusion

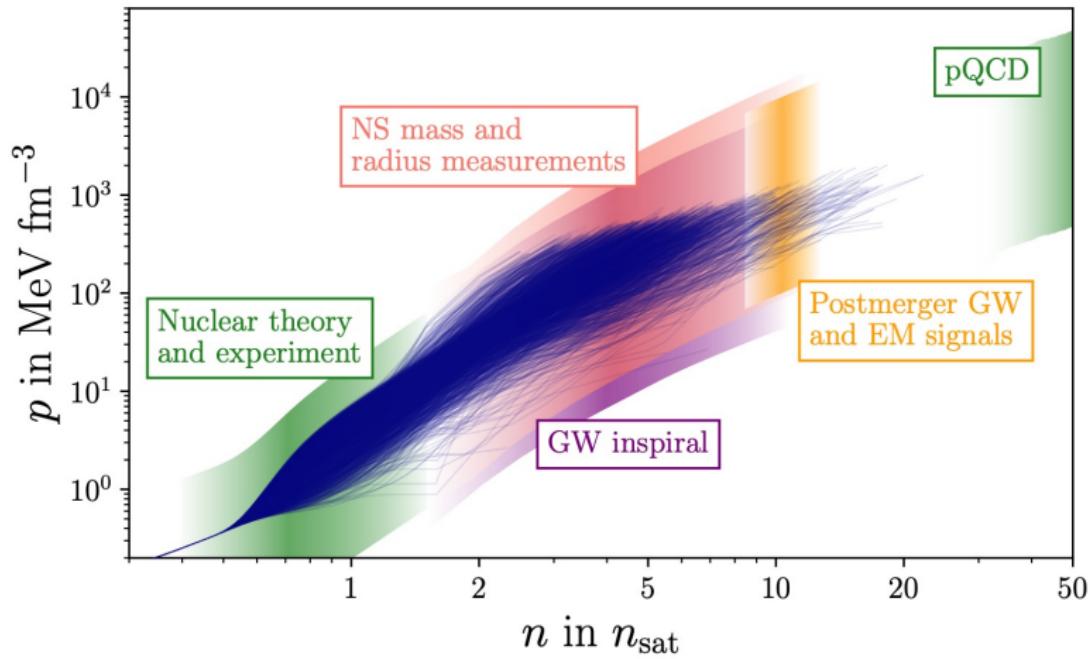
Neutron stars

- Neutron stars: supernova remnants, densest matter in the universe
- $m \sim 1.2 - 2.3 M_{\odot}$, $R \sim 10 - 13$ km



Equation of state

Neutron stars uniquely probe the equation of state (EOS) of dense nuclear matter [1]



Contents

① Introduction

Neutron stars and equation of state

Multimessenger astrophysics

Future detectors and challenges

② Methods

③ Applications

④ Neural priors

⑤ Outlook and conclusion

Multimessenger astrophysics

Observe neutron star mergers through gravitational waves and electromagnetic radiation: GW170817 [2, 3]

TODO: Snapshots

Contents

① Introduction

Neutron stars and equation of state

Multimessenger astrophysics

Future detectors and challenges

② Methods

③ Applications

④ Neural priors

⑤ Outlook and conclusion

Future GW detectors: Einstein Telescope

Einstein Telescope: Third-generation ground-based GW detector [4, 5]

Current software cannot handle the ET data analysis problem [6]

Future GW detectors: Einstein Telescope

Einstein Telescope: Third-generation ground-based GW detector [4, 5]

- Increased sensitivity
 - Louder signals
 - $10^5 - 10^6$ binary black hole mergers/year (now: $\sim 200/10$ years)
 - $10^4 - 10^5$ binary neutron star mergers/year (now: $2/10$ years)
 - $10^2 - 10^3$ multimessenger events/year (now: $1/10$ years)

Current software cannot handle the ET data analysis problem [6]

Future GW detectors: Einstein Telescope

Einstein Telescope: Third-generation ground-based GW detector [4, 5]

- Increased sensitivity
 - Louder signals
 - $10^5 - 10^6$ binary black hole mergers/year (now: $\sim 200/10$ years)
 - $10^4 - 10^5$ binary neutron star mergers/year (now: $2/10$ years)
 - $10^2 - 10^3$ multimessenger events/year (now: $1/10$ years)
- Observe from 5 Hz (now: 20 Hz)
 - Longer signals: binary neutron star: 2 hours vs 2 minutes
 - Signals will overlap

Current software cannot handle the ET data analysis problem [6]

My research focus: why – how – what

Why?

To make inference of multimessenger astrophysics scalable

- Prepare for future detectors
- Understand systematic effects through simulated data

My research focus: why – how – what

Why?

To make inference of multimessenger astrophysics scalable

- Prepare for future detectors
- Understand systematic effects through simulated data

How?

Without compromises to flexibility and accuracy

- Accelerate with GPU hardware, differentiable programming
- Use machine learning to assist inference, not replace it

My research focus: why – how – what

Why?

To make inference of multimessenger astrophysics scalable

- Prepare for future detectors
- Understand systematic effects through simulated data

How?

Without compromises to flexibility and accuracy

- Accelerate with GPU hardware, differentiable programming
- Use machine learning to assist inference, not replace it

What?

GPU-accelerated Bayesian joint inference framework of neutron star mergers and dense matter physics

Contents

① Introduction

② Methods

Bayesian inference and sampling

JAX

Normalizing flows

FLOWMC

③ Applications

④ Neural priors

⑤ Outlook and conclusion

Contents

① Introduction

② Methods

Bayesian inference and sampling

JAX

Normalizing flows

FLOWMC

③ Applications

④ Neural priors

⑤ Outlook and conclusion

Parameter estimation: Bayesian inference

How do we “measure” source parameters θ for data d ?

Bayesian inference:

$$\text{posterior} = p(\theta|d, \mathcal{M}) = \frac{p(d|\theta, \mathcal{M})p(\theta|\mathcal{M})}{p(d|\mathcal{M})} = \frac{\text{likelihood} \times \text{prior}}{\text{evidence}}$$

Ingredients:

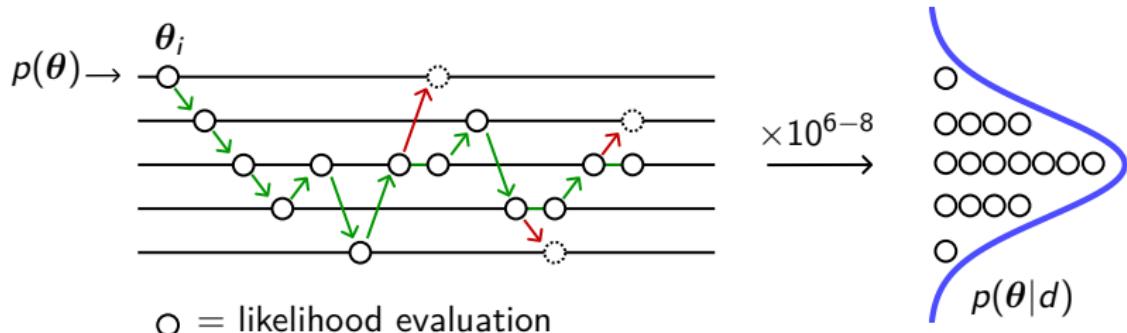
- Prior: specified by users
- Likelihood: often costly
- Posterior: intractable, needs stochastic samplers
- Evidence: model selection

Parameter estimation: Markov chain Monte Carlo

How do we sample the posterior? **Markov chain Monte Carlo**

- N chains θ_i explore posterior in parallel
- Evolve chains to new position: proposal
- Compute likelihood \rightarrow accept/reject

Alternatives: nested sampling, sequential Monte Carlo



Computational aspects

Total runtime $\approx N_{\text{likelihood}} \times \tau_{\text{likelihood}}$

Computational aspects

$$\text{Total runtime} \approx N_{\text{likelihood}} \times \tau_{\text{likelihood}}$$

- $N_{\text{likelihood}}$: total number of likelihood evaluations
 - Dimensionality of θ
 - Signal-to-noise ratio
 - Multimodality/shape posterior
 - Efficiency of sampler (proposals)

Computational aspects

$$\text{Total runtime} \approx N_{\text{likelihood}} \times \tau_{\text{likelihood}}$$

- $N_{\text{likelihood}}$: total number of likelihood evaluations
 - Dimensionality of θ
 - Signal-to-noise ratio
 - Multimodality/shape posterior
 - Efficiency of sampler (proposals)
- $\tau_{\text{likelihood}}$: speed of a single likelihood evaluation
 - Complexity model/likelihood function
 - Parallelization of likelihood evaluations
 - Hardware

Computational aspects

$$\text{Total runtime} \approx N_{\text{likelihood}} \times \tau_{\text{likelihood}}$$

- $N_{\text{likelihood}}$: total number of likelihood evaluations
 - Dimensionality of θ
 - Signal-to-noise ratio
 - Multimodality/shape posterior
 - **Efficiency of sampler (proposals)**
- $\tau_{\text{likelihood}}$: speed of a single likelihood evaluation
 - Complexity model/likelihood function
 - **Parallelization of likelihood evaluations**
 - **Hardware**

How can we optimize this?

Contents

① Introduction

② Methods

Bayesian inference and sampling

JAX

Normalizing flows

FLOWMC

③ Applications

④ Neural priors

⑤ Outlook and conclusion

jax “is a Python library for accelerator-oriented array computation and program transformation, designed for **high-performance** numerical computing and large-scale machine learning.” [7]



In particular:

- Python: integrate into existing workflows
- Focus on arrays: widespread applications
- Numpy-like API: (`numpy` → `jax.numpy`)

jax “is a Python library for accelerator-oriented array computation and program transformation, designed for **high-performance** numerical computing and large-scale machine learning.” [7]



In particular:

- Python: integrate into existing workflows
- Focus on arrays: widespread applications
- Numpy-like API: (numpy → jax.numpy)

Gamechangers:

- GPU accelerators
- Composable function transformations: `jit`, `grad`, `vmap` & `pmap`

JAX – Function transformations

- `jit`: Just-in-time compilation to XLA (CPU/GPU/TPU)

TODO: Some figure for autodiff

JAX – Function transformations

- `jit`: Just-in-time compilation to XLA (CPU/GPU/TPU)
- `grad`: Automatic differentiation: compute gradients with chain rule

TODO: Some figure for autodiff

JAX – Function transformations

- `jit`: Just-in-time compilation to XLA (CPU/GPU/TPU)
- `grad`: Automatic differentiation: compute gradients with chain rule
- `vmap`, `pmap`: Vectorization, batch processing, parallelization

TODO: Some figure for autodiff

JAX – Function transformations

- `jit`: Just-in-time compilation to XLA (CPU/GPU/TPU)
- `grad`: Automatic differentiation: compute gradients with chain rule
- `vmap`, `pmap`: Vectorization, batch processing, parallelization
- Composable, e.g.:
 - Higher-order derivatives: `jaxgrad(jaxgrad(f))`
 - Evolve N chains in parallel along gradient of the likelihood

$$\theta \leftarrow \alpha * \text{vmap}(\text{jit}(\text{grad}(\text{logL}(\theta))))$$

TODO: Some figure for autodiff

Computational aspects

$$\text{Total runtime} \approx N_{\text{likelihood}} \times \tau_{\text{likelihood}}$$

- $N_{\text{likelihood}}$: total number of likelihood evaluations
 - Dimensionality of θ
 - Signal-to-noise ratio
 - Multimodality/shape posterior
 - **Efficiency of sampler (proposals)**
- $\tau_{\text{likelihood}}$: speed of a single likelihood evaluation
 - Complexity model/likelihood function
 - **Parallelization of likelihood evaluations**
 - **Hardware**

How can we optimize this?

Contents

① Introduction

② Methods

Bayesian inference and sampling

JAX

Normalizing flows

FLOWMC

③ Applications

④ Neural priors

⑤ Outlook and conclusion

Normalizing flows

Q: Given samples $\{x\} \sim p(x)$, how can we get $p(x)$?

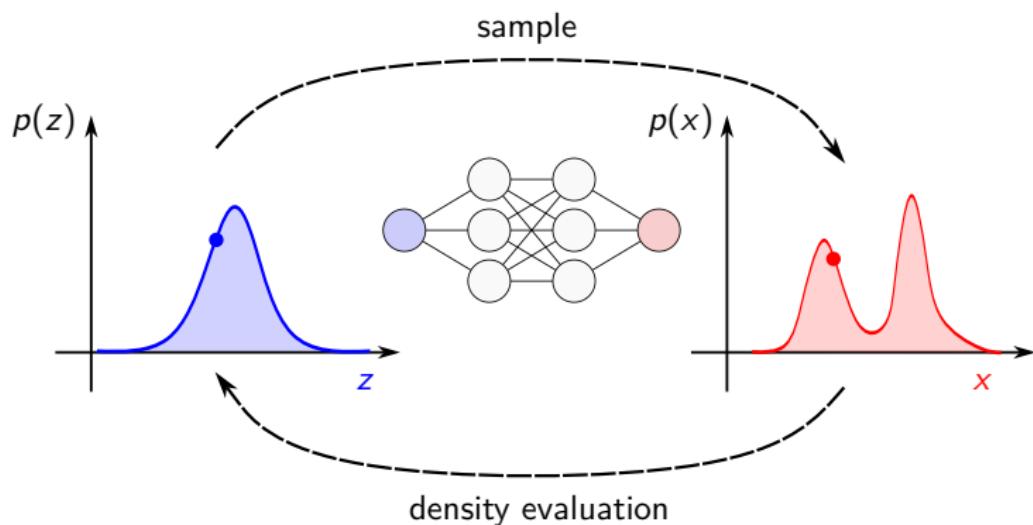
- Evaluate density at any point
- Generate new samples

Normalizing flows

Q: Given samples $\{x\} \sim p(x)$, how can we get $p(x)$?

- Evaluate density at any point
- Generate new samples

A: Normalizing flows: generative machine learning model, bijection between **latent** space (Gaussian) and **data** space



Contents

① Introduction

② Methods

Bayesian inference and sampling

JAX

Normalizing flows

FLOWMC

③ Applications

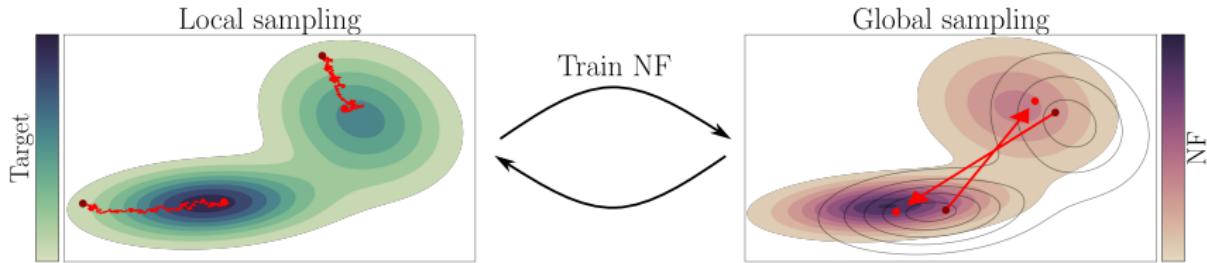
④ Neural priors

⑤ Outlook and conclusion

FLOWMC

FLOWMC  [8, 9]: normalizing flow-enhanced MCMC

- ① Run gradient-based sampler (local sampler)
- ② Train normalizing flow on MCMC chains
- ③ Propose samples from the normalizing flow (global sampler)
- ④ Repeat



Contents

① Introduction

② Methods

③ Applications

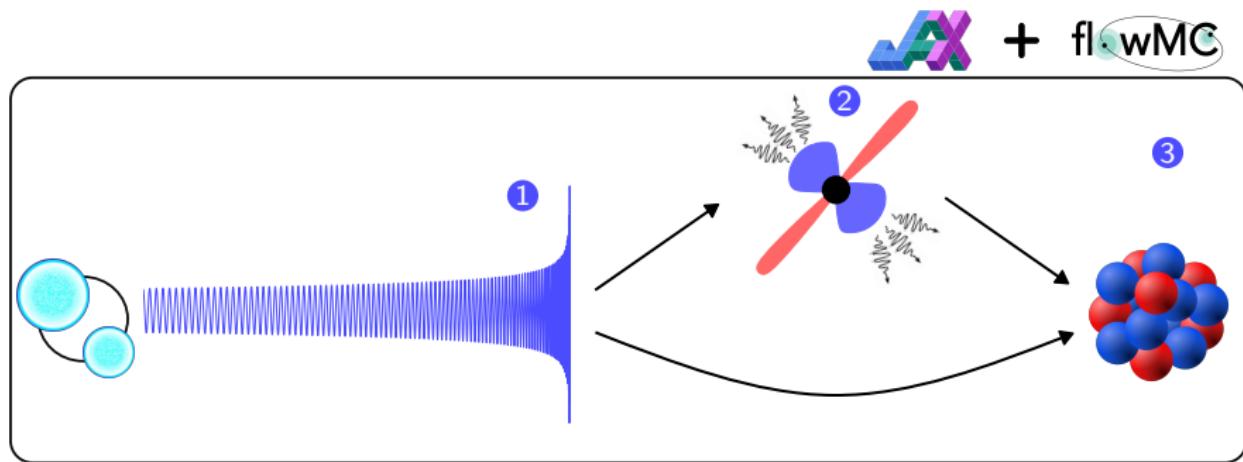
④ Neural priors

⑤ Outlook and conclusion

Overview

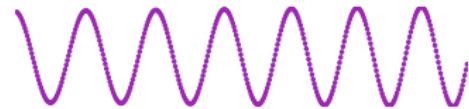
Analyzing a multi-messenger **binary neutron star** signal:

- ① Gravitational waves
- ② Electromagnetic counterparts
- ③ Nuclear equation of state



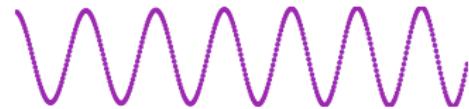
Gravitational waves: RIPPLE

- Waveforms on GPU: $\mathcal{O}(10^3)$ faster
- RIPPLE [10]: from LALSUITE to JAX



Gravitational waves: RIPPLE

- Waveforms on GPU: $\mathcal{O}(10^3)$ faster
- RIPPLE [10]: from LALSUITE to JAX
- Waveforms available in RIPPLE:
 - Binary black holes:
 - IMRPhenomXAS
 - IMRPhenomD
 - IMRPhenomPv2
 - Binary neutron stars:
 - TaylorF2
 - IMRPhenomD_NRTidalv2
 - IMRPhenomPv_NRTidalv2 (by Nihar Gupte)
- Also see GWFEST [11] and SFTS [12]



Gravitational waves: JIM

Parameter estimation: JIM Ω [13, 14]: RIPPLE + FLOWMC

- Verified for binary neutron stars in current generation detectors [14]
- Matches BILBY, passes *pp*-test

Gravitational waves: JIM

Parameter estimation: JIM  [13, 14]: RIPPLE + FLOWMC

- Verified for binary neutron stars in current generation detectors [14]
- Matches BILBY, passes *pp*-test
- Runtime: from hours to minutes

Event	Waveform	JIM (1 GPU)	pBILBY (480 cores)	RB-BILBY (24 cores)	ROQ-BILBY (24 cores)
GW170817	TF2	15.33 min	9.64 h	3.8 h	–
	NRTv2	25.59 min	10.99 h	4.11 h	1.65 h
GW190425	TF2	18.30 min	8.18 h	2.81 h	–
	NRTv2	21.20 min	4.91 h	2.42 h	0.97 h
Injection	TF2	24.76 min	–	–	–
	NRTv2	18.02 min	–	–	–

Gravitational waves: JIM

Parameter estimation: JIM  [13, 14]: RIPPLE + FLOWMC

- Verified for binary neutron stars in current generation detectors [14]
- Matches BILBY, passes *pp*-test
- Runtime: from hours to minutes
- Up to $5 - 10 \times$ more effective samples (normalizing flow)

Effective sample size	
JIM	4.1×10^4
PBILBY	6.5×10^3
RB-BILBY	5.8×10^3
ROQ-BILBY	7.4×10^3

Gravitational waves: JIM

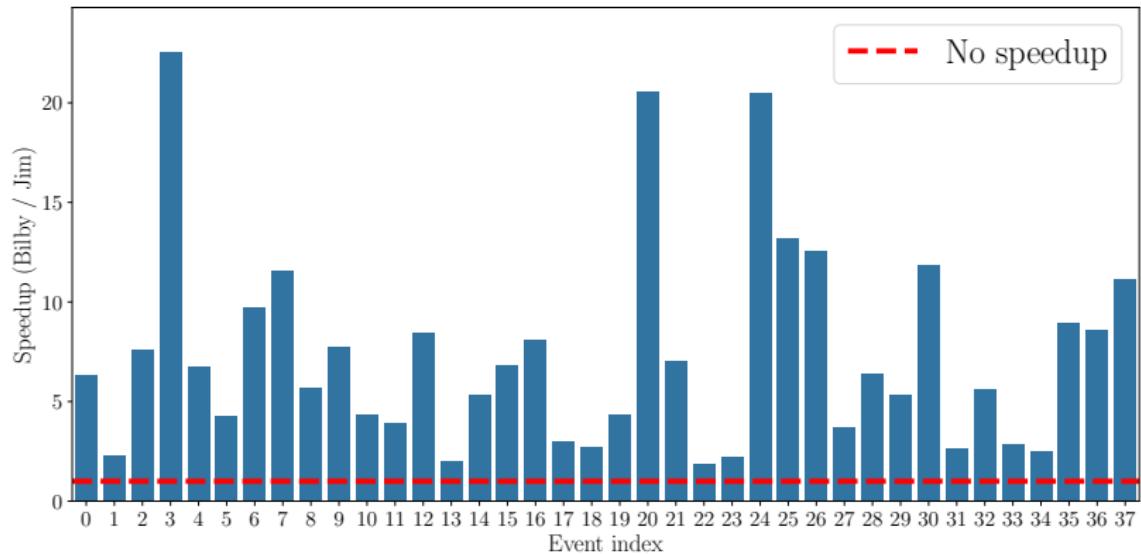
Parameter estimation: JIM  [13, 14]: RIPPLE + FLOWMC

- Verified for binary neutron stars in current generation detectors [14]
- Matches BILBY, passes *pp*-test
- Runtime: from hours to minutes
- Up to 5 – 10× more effective samples (normalizing flow)
- More cost-effective/energy-efficient

	kWh	CO ₂ [10 ³ kg]	Trees
JIM	34	11	0.55
PBILBY	4127	1354	67.68
RB-BILBY	80	26	1.32
ROQ-BILBY	sampling	32	0.52
	precompute	27	0.44

Gravitational waves: GWTC-3 analysis – Thomas Ng

- Re-analyze GWTC-3 with `IMRPhenomPv2`
- JIM on A800 GPU: on average ~ 7 times faster than BILBY on 16 CPUs, **without compromises**
- **Work in progress!**

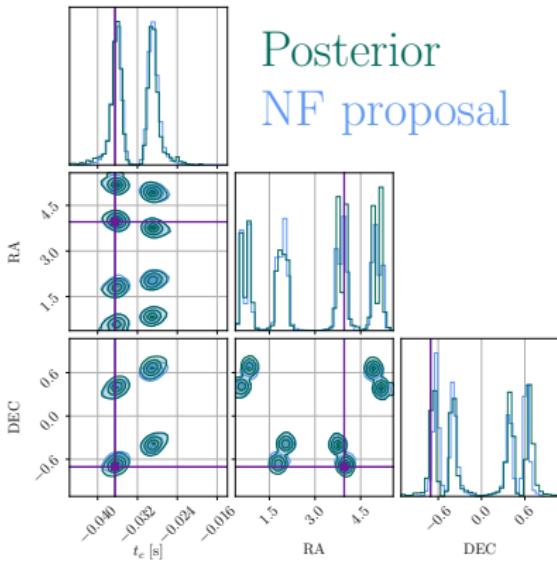


Work in progress: showing proof-of-concepts!

Einstein Telescope

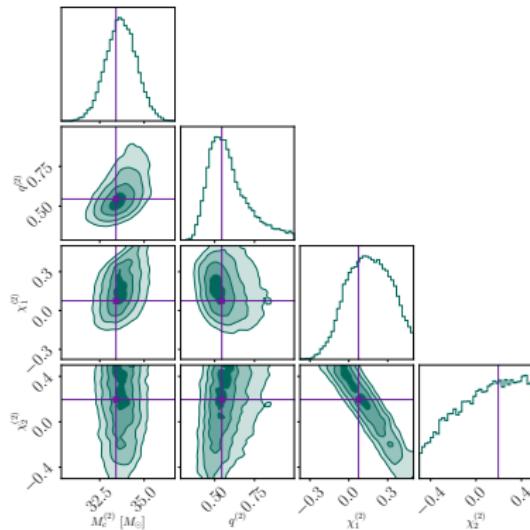
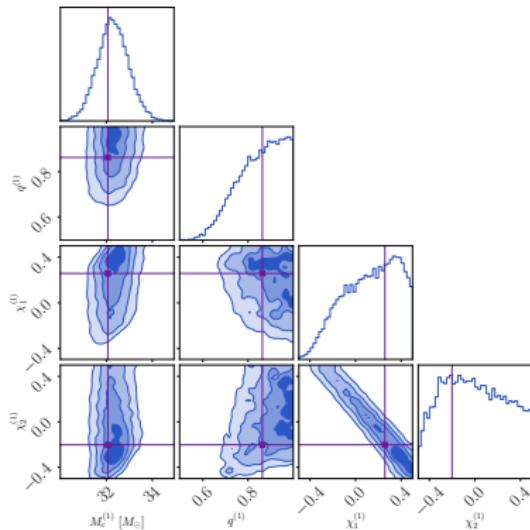
Work in progress: showing proof-of-concepts!

- ET posteriors are multimodal
- Normalizing flows help to jump between modes



Overlapping signals

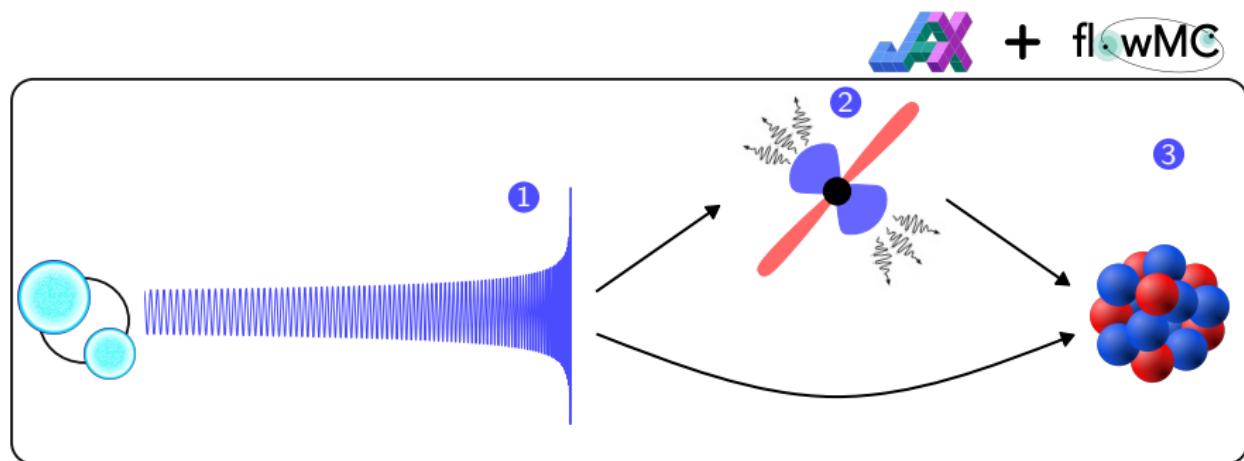
- Assess scaling of JIM: BBH+BBH with LIGO-Virgo
 - 2 binary black hole mergers: 22 parameters
 - $M_c^{(1)} = 32M_\odot$, $M_c^{(2)} = 33M_\odot$, $\Delta t = 70$ ms
 - $\text{SNR}^{(1)} = 25.76$, $\text{SNR}^{(2)} = 25.24$
 - 1h28m on H100 GPU (vs several days on 16 CPUs [15])



Overview

Analyzing a multi-messenger **binary neutron star** signal:

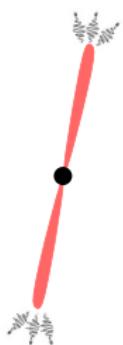
- ① Gravitational waves
- ② **Electromagnetic counterparts**
- ③ Nuclear equation of state



- Kilonovae and **GRB afterglows**
- Code libraries too large to 'jaxify' (e.g. AFTERGLOWPY [16])

Electromagnetic counterparts (Hauke Koehn, Tim Dietrich)

- Kilonovae and **GRB afterglows**
- Code libraries too large to 'jaxify' (e.g. AFTERGLOWPY [16])
- Neural network emulators for inference: FIESTA  [17]

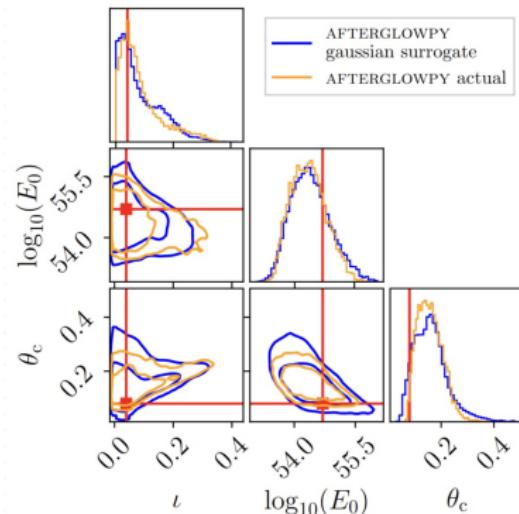


FIESTA

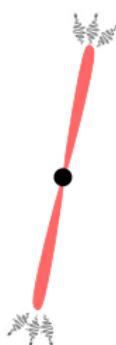
- 1m36s
- 1 H100 GPU

AFTERGLOWPY

- 4 hours
- 30 CPUs



- Kilonovae and **GRB afterglows**
- Code libraries too large to ‘jaxify’ (e.g. AFTERGLOWPY [16])
- Neural network emulators for inference: FIESTA  [17]
- Scales well for systematics ‘nuisance parameters’

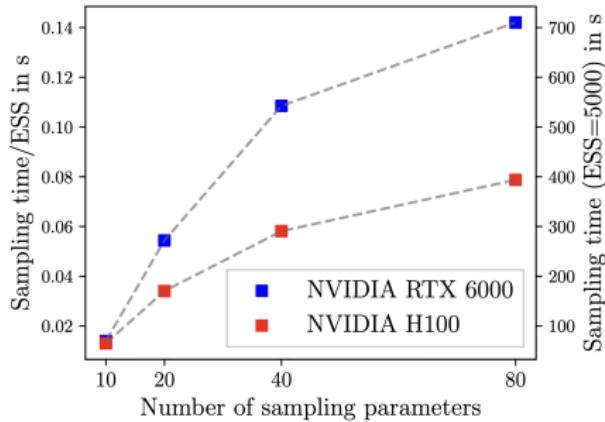


FIESTA

- 1m36s
- 1 H100 GPU

AFTERGLOWPY

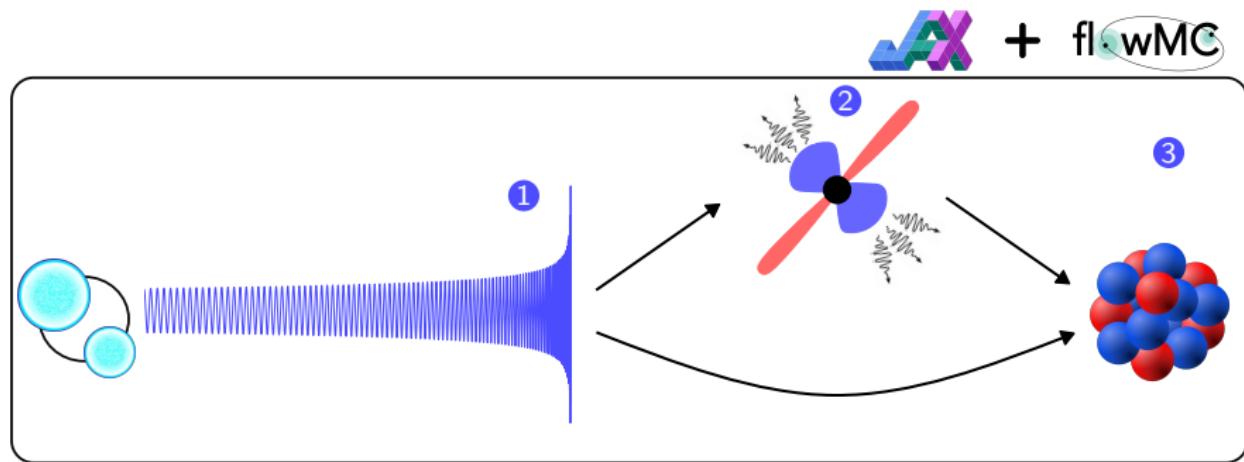
- 4 hours
- 30 CPUs



Overview

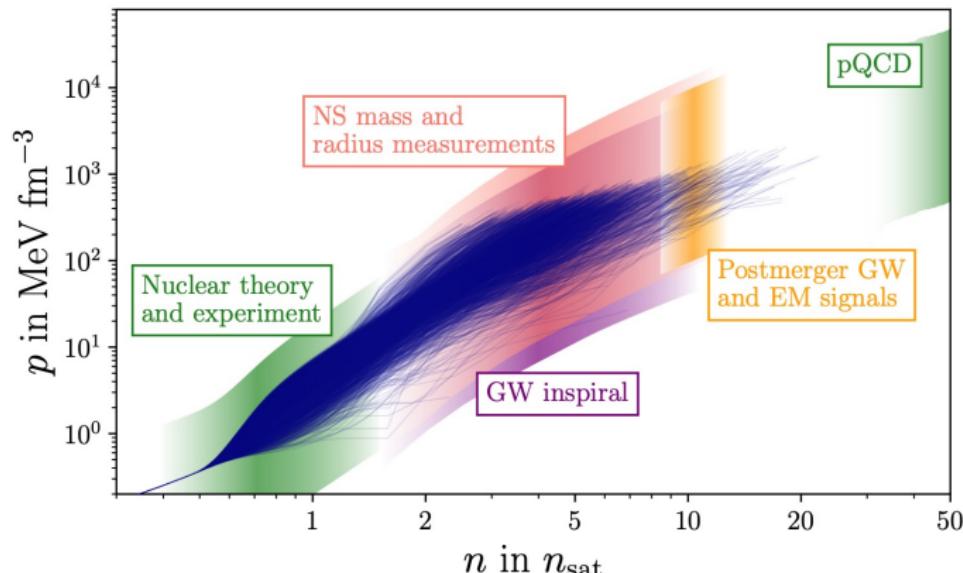
Analyzing a multi-messenger **binary neutron star** signal:

- ① Gravitational waves
- ② Electromagnetic counterparts
- ③ Nuclear equation of state



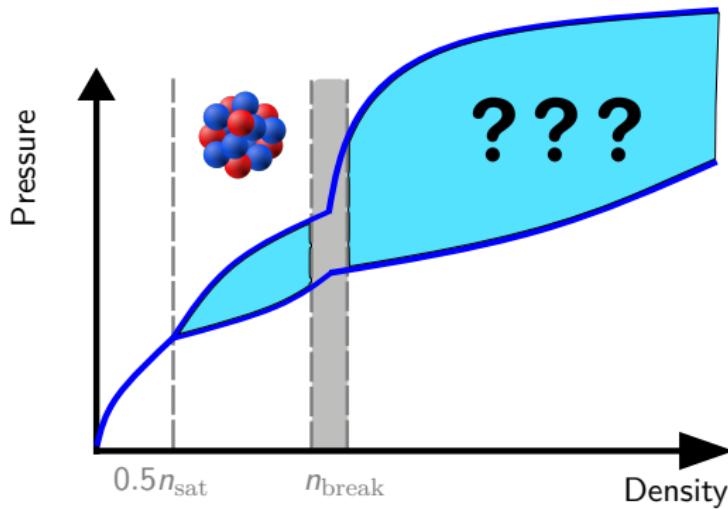
Equation of state inference – warmup

- How do we constrain the EOS from neutron star observations?
- Define parametrization and likelihood
- GWs from binary neutron stars inspiral as example



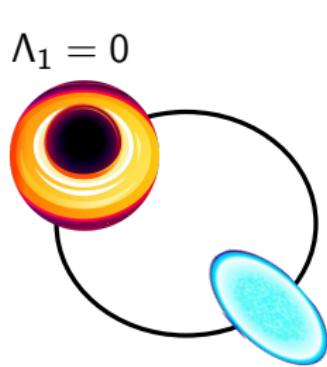
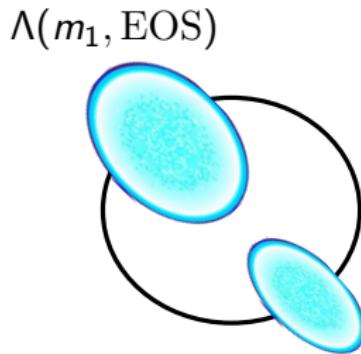
Equation of state inference – parametrization

- $n \leq \frac{1}{2}n_{\text{sat}}$: fixed
- $\frac{1}{2}n_{\text{sat}} \leq n \leq n_{\text{break}}$: metamodel EOS, nuclear physics inspired
- $n \leq n_{\text{break}}$: agnostic extension
- > 26 (!) parameters θ_{EOS}



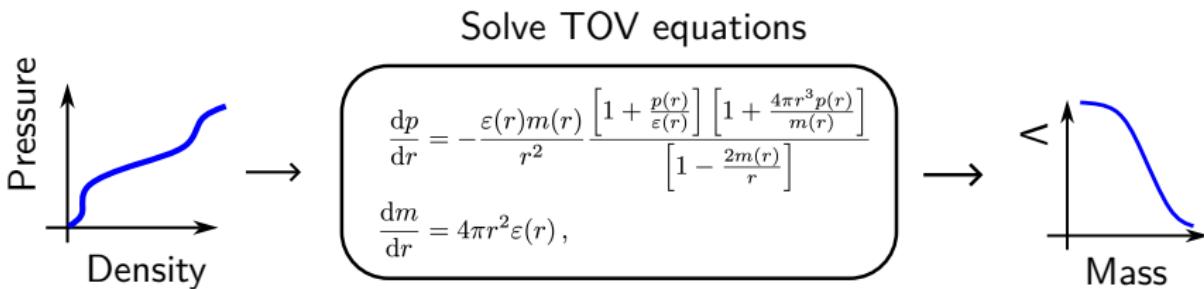
Tidal deformability

- Neutron stars are tidally deformed in a binary, quantified by tidal deformability Λ : affect phase of GWs
- Neutron stars: $\Lambda = \Lambda(m, \text{EOS})$, black holes: $\Lambda = 0$
- GWs give us $p(m_1, m_2, \Lambda_1, \Lambda_2 | d)$ for EOS inference


$$\Lambda(m_2, \text{EOS})$$

$$\Lambda(m_2, \text{EOS})$$

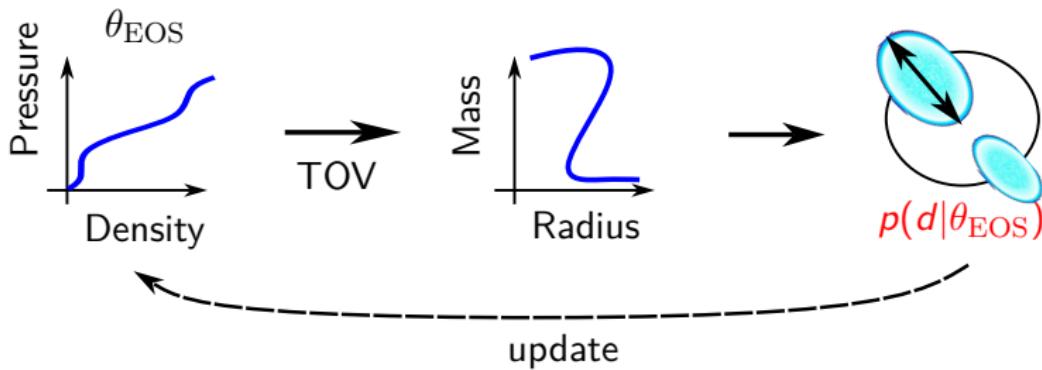
Equation of state

- To predict neutron star properties, we solve the TOV equations: ordinary differential equations (ODEs)



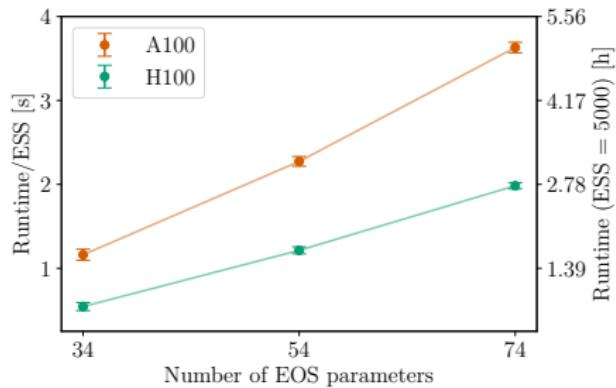
Equation of state

- To predict neutron star properties, we solve the TOV equations: ordinary differential equations (ODEs)
- Done for each sample θ_{EOS} : **costly likelihood**
- How to make this scalable without compromises?

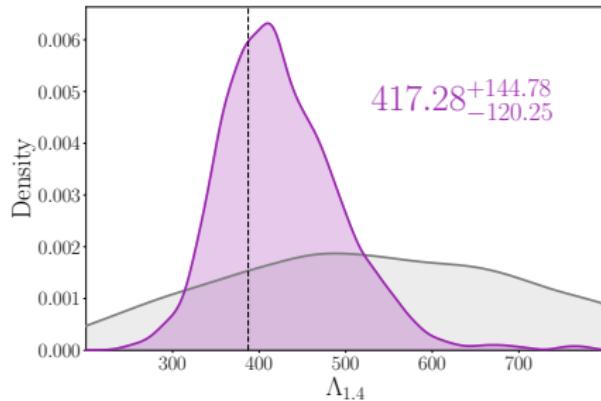
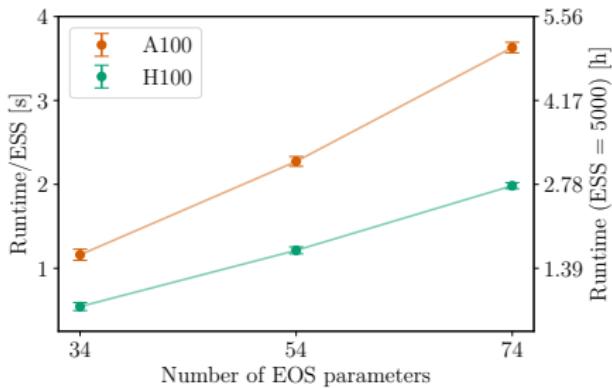


JESTER

- JESTER  [18]: JAX-based TOV solver
 - $1000\times$ faster, without compromises!
 - Full inference in \sim hours



- JESTER  [18]: JAX-based TOV solver
 - $1000\times$ faster, without compromises!
 - Full inference in \sim hours
- JIM+JESTER: from GWs to EOS in a few hours
 - Example: 20 binary neutron stars in O5
- Open up possibilities to study systematics in EOS inference



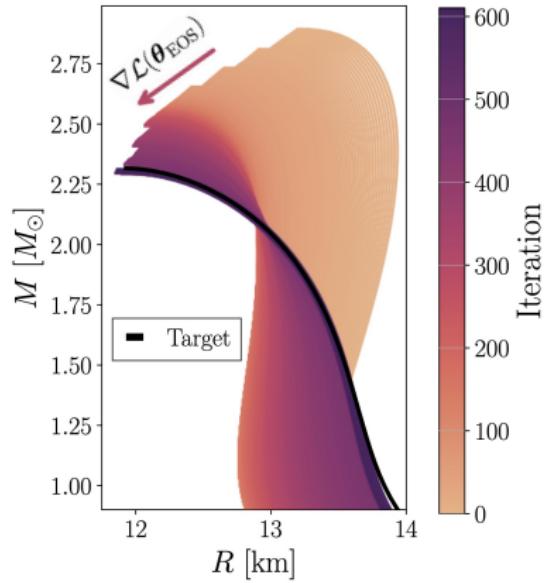
Auto-differentiable ODE solvers

- ODE solvers written in JAX are auto-differentiable
- Frame inference as optimization problem:
 - Gradient descent on loss function $\mathcal{L}(\theta_{\text{EOS}})$

Auto-differentiable ODE solvers

- ODE solvers written in JAX are auto-differentiable
- Frame inference as optimization problem:
 - Gradient descent on loss function $\mathcal{L}(\theta_{\text{EOS}})$

$$\mathcal{L}(\theta_{\text{EOS}}) = \frac{1}{N} \sum_{i=1}^N \left| \frac{R_i(\theta_{\text{EOS}}) - \hat{R}_i}{\hat{R}_i} \right|$$



Contents

① Introduction

② Methods

③ Applications

④ Neural priors

⑤ Outlook and conclusion

and now, for something completely different...

Contents

① Introduction

② Methods

③ Applications

④ Neural priors

⑤ Outlook and conclusion

Conclusion

- Progress on scalable Bayesian inference, with minimal pre-training
- Hybrid acceleration: GPUs + normalizing flow proposals
 - JAX/GPU: faster likelihoods
 - FLOWMC: sampling converges faster
- Simulators in JAX can remove the need for emulators (GW, TOV)
- Auto-differentiable ODE solvers: inference as optimization problem

Let's talk!



Thanks for listening!

References I

- [1] Hauke Koehn et al. "From existing and new nuclear and astrophysical constraints to stringent limits on the equation of state of neutron-rich dense matter". In: (Feb. 2024). arXiv: [2402.04172 \[astro-ph.HE\]](https://arxiv.org/abs/2402.04172).
- [2] B. P. Abbott et al. "GW170817: Observation of Gravitational Waves from a Binary Neutron Star Inspiral". In: *Phys. Rev. Lett.* 119.16 (2017), p. 161101. DOI: [10.1103/PhysRevLett.119.161101](https://doi.org/10.1103/PhysRevLett.119.161101). arXiv: [1710.05832 \[gr-qc\]](https://arxiv.org/abs/1710.05832).
- [3] B. P. Abbott et al. "Gravitational Waves and Gamma-rays from a Binary Neutron Star Merger: GW170817 and GRB 170817A". In: *Astrophys. J. Lett.* 848.2 (2017), p. L13. DOI: [10.3847/2041-8213/aa920c](https://doi.org/10.3847/2041-8213/aa920c). arXiv: [1710.05834 \[astro-ph.HE\]](https://arxiv.org/abs/1710.05834).
- [4] Michele Maggiore et al. "Science Case for the Einstein Telescope". In: *JCAP* 03 (2020), p. 050. DOI: [10.1088/1475-7516/2020/03/050](https://doi.org/10.1088/1475-7516/2020/03/050). arXiv: [1912.02622 \[astro-ph.CO\]](https://arxiv.org/abs/1912.02622).
- [5] Adrian Abac et al. "The Science of the Einstein Telescope". In: (Mar. 2025). arXiv: [2503.12263 \[gr-qc\]](https://arxiv.org/abs/2503.12263).
- [6] Qian Hu and John Veitch. "Costs of Bayesian Parameter Estimation in Third-Generation Gravitational Wave Detectors: a Review of Acceleration Methods". In: (Dec. 2024). arXiv: [2412.02651 \[gr-qc\]](https://arxiv.org/abs/2412.02651).
- [7] James Bradbury et al. *JAX: composable transformations of Python+NumPy programs*. Version 0.3.13. 2018. URL: <http://github.com/jax-ml/jax>.

References II

- [8] Marylou Gabrié, Grant M. Rotskoff, and Eric Vanden-Eijnden. “Adaptive Monte Carlo augmented with normalizing flows”. In: *Proc. Nat. Acad. Sci.* 119.10 (2022), e2109420119. DOI: [10.1073/pnas.2109420119](https://doi.org/10.1073/pnas.2109420119). arXiv: [2105.12603 \[physics.data-an\]](https://arxiv.org/abs/2105.12603).
- [9] Kaze W. k. Wong, Marylou Gabrié, and Daniel Foreman-Mackey. “flowMC: Normalizing flow enhanced sampling package for probabilistic inference in JAX”. In: *J. Open Source Softw.* 8.83 (2023), p. 5021. DOI: [10.21105/joss.05021](https://doi.org/10.21105/joss.05021). arXiv: [2211.06397 \[astro-ph.IM\]](https://arxiv.org/abs/2211.06397).
- [10] Thomas D. P. Edwards et al. “Differentiable and hardware-accelerated waveforms for gravitational wave data analysis”. In: *Phys. Rev. D* 110.6 (2024), p. 064028. DOI: [10.1103/PhysRevD.110.064028](https://doi.org/10.1103/PhysRevD.110.064028). arXiv: [2302.05329 \[astro-ph.IM\]](https://arxiv.org/abs/2302.05329).
- [11] Francesco Iacovelli et al. “GWFEST: A Fisher Information Matrix Python Code for Third-generation Gravitational-wave Detectors”. In: *Astrophys. J. Supp.* 263.1 (2022), p. 2. DOI: [10.3847/1538-4365/ac9129](https://doi.org/10.3847/1538-4365/ac9129). arXiv: [2207.06910 \[astro-ph.IM\]](https://arxiv.org/abs/2207.06910).
- [12] Rodrigo Tenorio and Davide Gerosa. “Scalable data-analysis framework for long-duration gravitational waves from compact binaries using short Fourier transforms”. In: *Phys. Rev. D* 111.10 (2025), p. 104044. DOI: [10.1103/PhysRevD.111.104044](https://doi.org/10.1103/PhysRevD.111.104044). arXiv: [2502.11823 \[gr-qc\]](https://arxiv.org/abs/2502.11823).

References III

- [13] Kaze W. K. Wong, Maximiliano Isi, and Thomas D. P. Edwards. “Fast Gravitational-wave Parameter Estimation without Compromises”. In: *Astrophys. J.* 958.2 (2023), p. 129. DOI: [10.3847/1538-4357/acf5cd](https://doi.org/10.3847/1538-4357/acf5cd). arXiv: [2302.05333](https://arxiv.org/abs/2302.05333) [astro-ph.IM].
- [14] Thibeau Wouters et al. “Robust parameter estimation within minutes on gravitational wave signals from binary neutron star inspirals”. In: *Phys. Rev. D* 110.8 (2024), p. 083033. DOI: [10.1103/PhysRevD.110.083033](https://doi.org/10.1103/PhysRevD.110.083033). arXiv: [2404.11397](https://arxiv.org/abs/2404.11397) [astro-ph.IM].
- [15] Justin Januart et al. “Analyses of overlapping gravitational wave signals using hierarchical subtraction and joint parameter estimation”. In: *Mon. Not. Roy. Astron. Soc.* 523.2 (2023), pp. 1699–1710. DOI: [10.1093/mnras/stad1542](https://doi.org/10.1093/mnras/stad1542). arXiv: [2211.01304](https://arxiv.org/abs/2211.01304) [gr-qc].
- [16] Geoffrey Ryan et al. “Gamma-Ray Burst Afterglows in the Multimessenger Era: Numerical Models and Closure Relations”. In: *Astrophys. J.* 896.2 (2020), p. 166. DOI: [10.3847/1538-4357/ab93cf](https://doi.org/10.3847/1538-4357/ab93cf). arXiv: [1909.11691](https://arxiv.org/abs/1909.11691) [astro-ph.HE].
- [17] Hauke Koehn et al. “Efficient Bayesian analysis of kilonovae and GRB afterglows with fiesta”. In: (July 2025). arXiv: [2507.13807](https://arxiv.org/abs/2507.13807) [astro-ph.HE].
- [18] Thibeau Wouters et al. “Leveraging differentiable programming in the inverse problem of neutron stars”. In: (Apr. 2025). arXiv: [2504.15893](https://arxiv.org/abs/2504.15893) [astro-ph.HE].

References IV

- [19] Kurzgesagt. *Figures taken from “Neutron Stars - The Most Extreme Things that are not Black Holes”*. Accessed on May 14, 2025. 2019. URL: <https://www.youtube.com/watch?v=udFxKZRyQt4>.
- [20] Hergé. *Cover figure created with ChatGPT using this input figure from the comic Destination Moon*. Accessed on May 14, 2025. 2019. URL: <https://www.youtube.com/watch?v=udFxKZRyQt4>.
- [21] Jason D. McEwen et al. *Machine learning assisted Bayesian model comparison: learnt harmonic mean estimator*. 2023. arXiv: 2111.12720 [stat.ME]. URL: <https://arxiv.org/abs/2111.12720>.
- [22] Alicja Polanska et al. *Learned harmonic mean estimation of the marginal likelihood with normalizing flows*. 2024. arXiv: 2307.00048 [stat.ME]. URL: <https://arxiv.org/abs/2307.00048>.
- [23] Alicja Polanska et al. “Accelerated Bayesian parameter estimation and model selection for gravitational waves with normalizing flows”. In: *38th conference on Neural Information Processing Systems*. Oct. 2024. arXiv: 2410.21076 [astro-ph.IM].

Evidence calculation: HARMONIC I

Evidence Z can be computed from posterior samples with HARMONIC [21] with the **harmonic mean estimator**

$$\begin{aligned}\rho &\equiv \mathbb{E}_{P(\theta|d)} \left[\frac{1}{L(\theta)} \right] \\ &= \int d\theta \frac{1}{L(\theta)} P(\theta|d) \\ &= \int d\theta \frac{1}{L(\theta)} \frac{\mathcal{L}(\theta)\pi(\theta)}{Z} = \frac{1}{Z}\end{aligned}$$

Therefore, estimate ρ with posterior samples:

$$\hat{\rho} = \frac{1}{N} \sum_{i=1}^N \frac{1}{L(\theta_i)}, \quad \theta_i \sim P(\theta|d)$$

Evidence calculation: HARMONIC II

Can be interpreted as importance sampling

$$\rho = \int d\theta \frac{1}{Z} \frac{\pi(\theta)}{P(\theta|d)} P(\theta|d),$$

but with target = prior and sampling density = posterior. Therefore, importance sampling is inefficient – how to solve?

New proposal:

$$\begin{aligned}\rho &= \mathbb{E}_{P(\theta|d)} \left[\frac{\varphi(\theta)}{\mathcal{L}(\theta)\pi(\theta)} \right] \\ &= \int d\theta \frac{\varphi(\theta)}{\mathcal{L}(\theta)\pi(\theta)} P(\theta|d) \\ &= \int d\theta \frac{\varphi(\theta)}{\mathcal{L}(\theta)\pi(\theta)} \frac{\mathcal{L}(\theta)\pi(\theta)}{Z} = \frac{1}{Z}\end{aligned}$$

Evidence calculation: HARMONIC III

Use the following estimator:

$$\hat{\rho} = \frac{1}{N} \sum_{i=1}^N \frac{\varphi(\theta_i)}{\mathcal{L}(\theta_i)\pi(\theta_i)}, \quad \theta_i \sim P(\theta|d)$$

Replace the target distribution π with φ : only requirement is that it is normalized

In practice, this can be achieved with a normalizing flow [22].

This has been verified to give accurate evidences (similar values as nested sampling) when GW posteriors are used [23].

HARMONIC with JIM [23]

Table 1: Total wall times to compute the evidence estimates for the examples discussed in the main text. We run BILBY on 16 CPU cores and JIM + harmonic on 1 GPU.

Example	Method	$\log(z)$	Sampling time	Evidence estimation time
4D	BILBY	390.33 ± 0.11	31.3 min	—
	JIM + harmonic	$390.360^{+0.006}_{-0.006}$	3.4 min	1.9 min
11D	BILBY	378.29 ± 0.15	3.5 h	—
	JIM + harmonic	$378.420^{+0.09}_{-0.08}$	11.8 min	2.4 min

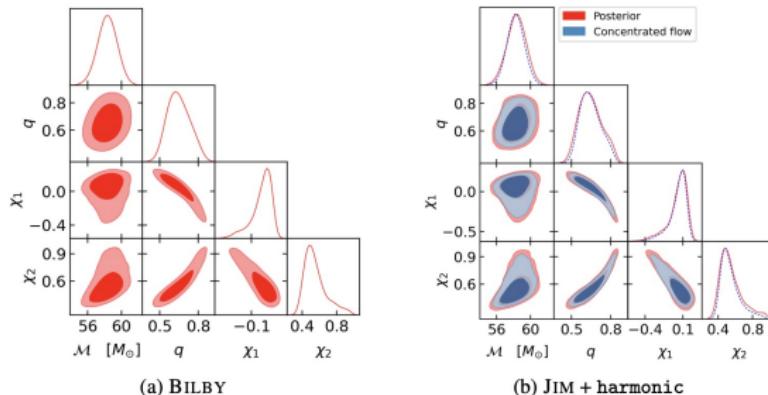
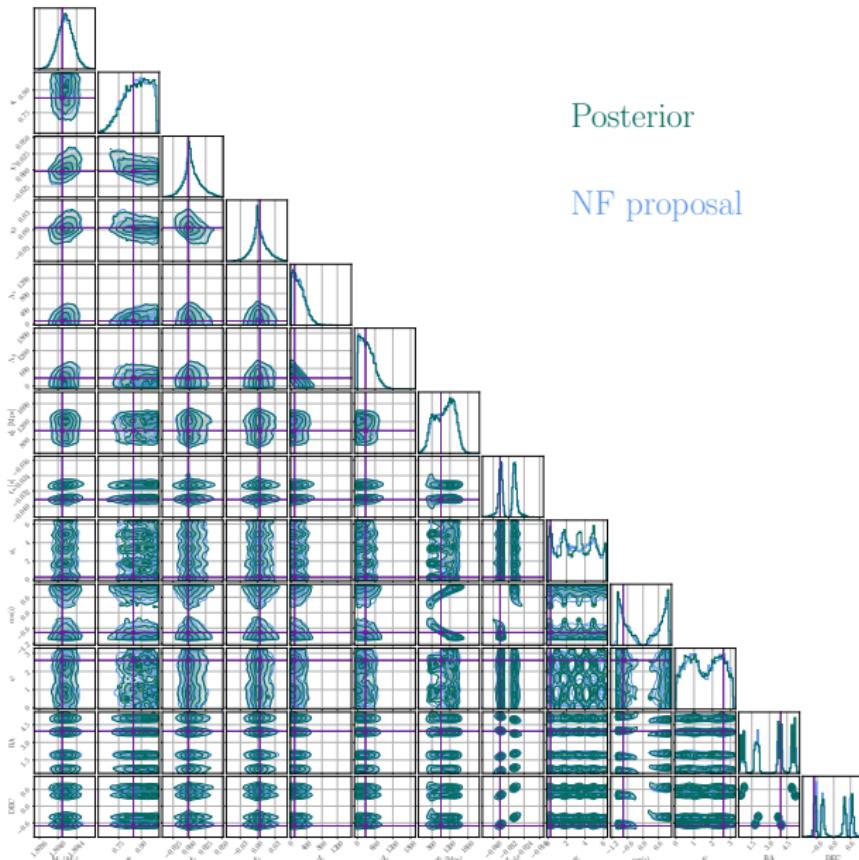
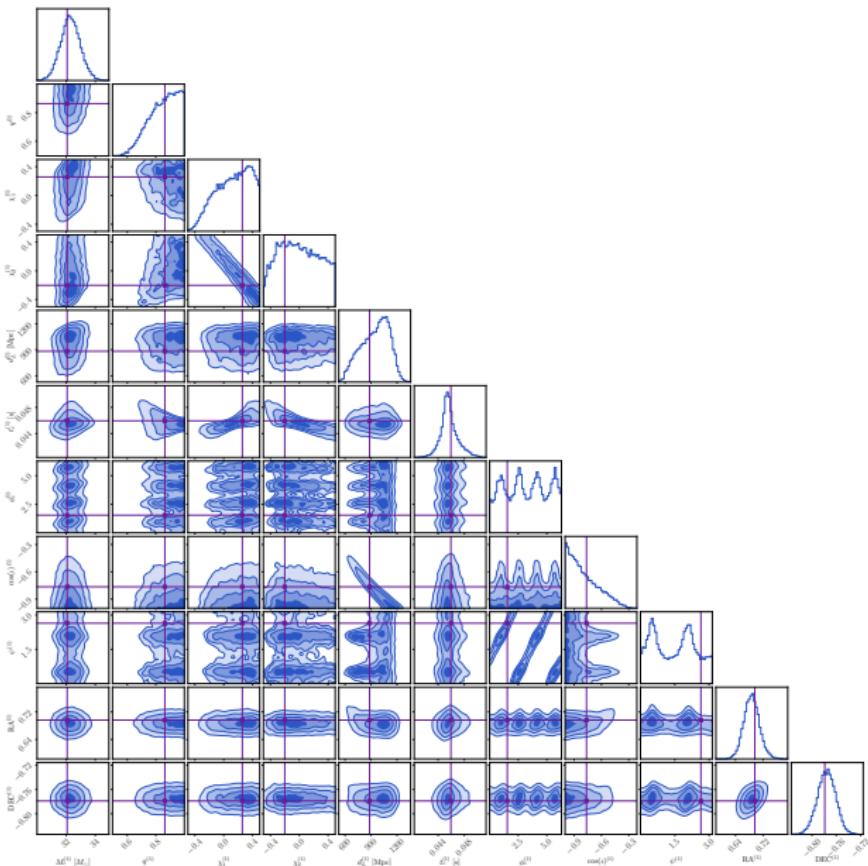


Figure 1: Corner plots for the 4-dimensional posterior samples from (a) BILBY and (b) JIM used for inference (solid red) alongside the concentrated flow at $T = 0.8$ used in the learned harmonic mean (dashed blue).

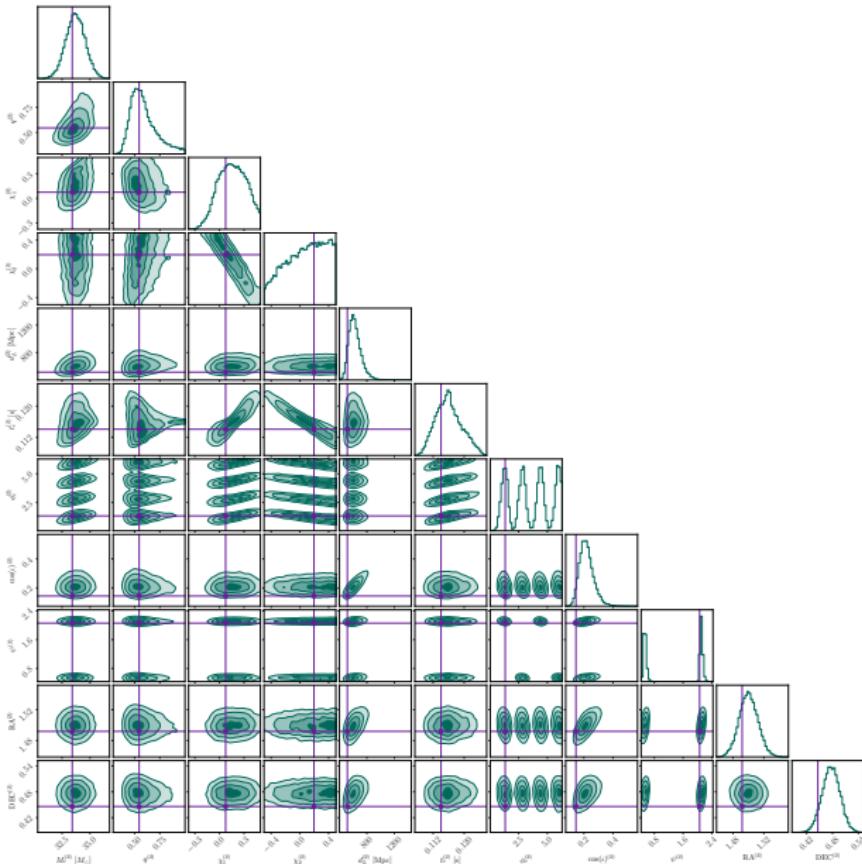
BNS in ET- Δ example: all parameters



Overlapping signals: all parameters signal A

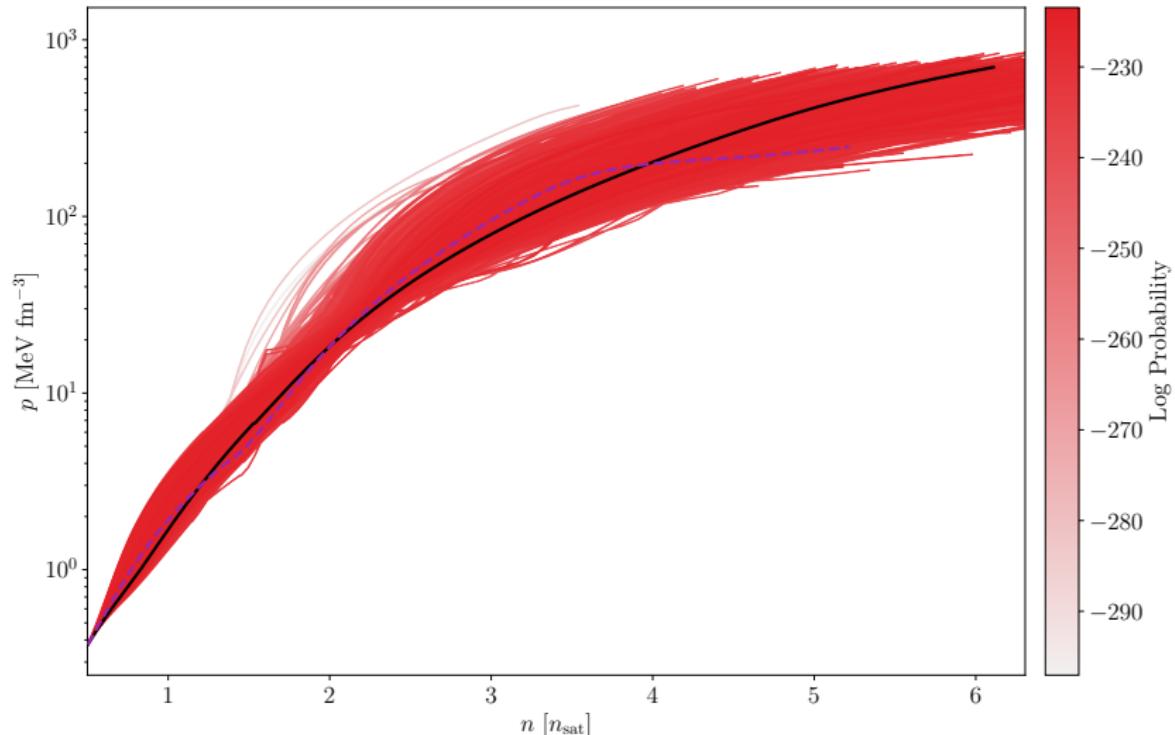


Overlapping signals: all parameters signal B



Equation of state O5 projection with 20 BNS: EOS

- **Purple:** target
- **Red:** posterior EOS samples (**black**: maximum log posterior)



Equation of state O5 projection with 20 BNS: NS

